

Infra-as-code is slow to write — and **dangerous to apply**

- ✗ Writing correct Terraform needs **specialist knowledge** — most people can't.
- ✗ Hand-written HCL is **error-prone**: wrong syntax, insecure defaults, drift.
- ✗ A single wrong **terraform apply** can **break prod** or **open a security hole** (public bucket, open port).
- ✗ Reviews are a **bottleneck** — plans are hard to read and slow to sign off.

The gap: teams want the speed of "just describe it" **without** giving up review, security, and a human approval before anything changes.

Describe it in English → reviewed Terraform

An **IaC Generator Agent**: you type what you want in plain language; the agent writes secure Terraform, runs a real **terraform plan**, explains the diff, checks it against policy — and applies *only after a human approves*.

01

Generate

gpt-4o writes secure-by-default HCL (AWS v5).

02

Plan

Real **terraform plan** — see exactly what changes. Nothing applied.

03

Explain

The agent reviews the diff in plain English + risks.


04

Govern

Policy check + human-gated apply.

The agent proposes; a human disposes. The LLM never touches your cloud — Terraform does, and only the read-only plan runs on its own.

The **process**, end to end



Connects to: OpenAI (gpt-4o) · the Terraform CLI · your AWS account (plan authenticates read-only). First `terraform init` caches the provider; every plan after is fast. Runs as a Streamlit app — `streamlit run app.py`.